

# SCALABM

DAVID R. PUGH, DANIEL F. TANG, J. DOYNE FARMER

## 1. OBJECTIVE

Our objective is to create a *user-friendly* toolkit for building *scalable*, *data-driven* agent-based models (ABMs) of *economic* systems.

### 1.1. Requirements.

1.1.1. ...*user-friendly*... We expect that users of our ABM toolkit can be classified as either consumers or producers. Consumers are users whose primary objective is to learn about the mechanisms driving an existing model's key results by playing with model parameters and components. Producers are users whose primary objective is to develop new models using some combination of existing and novel model components.

We want consumers of models built using our framework, particularly those not directly involved with development of any particular model, to be able to easily interact with a model in order to develop intuition and understanding about the mechanisms driving that model's key results.

- Models built using our toolkit should have browser-based user interfaces (UIs). Implementing browser-based UIs would allow us to leverage existing, high-quality Javascript libraries for real-time data visualization.
- Users should be able to “play-with” the model (i.e., tweak model parameters, substitute model components, etc.) from the browser in real-time (i.e., whilst a model is running).

We also want to develop a framework that minimizes developer time when either building a new model (or re-configuring an existing model).

- Models built using our toolkit should be composed of mostly existing components. This reduces development time for a new model to that needed to create a few novel components together with the time needed to wire all the desired model components together.
- The process of wiring model components together, sometimes called dependency injection (DI), should be as simple and transparent as possible.<sup>1</sup>

---

*Date:* September 27, 2015.

<sup>1</sup> There are many popular DI approaches (i.e., Guice, Spring, etc). While our framework should not depend on any particular DI library, we should think carefully about what DI library we choose as I expect many users will just mimic our choice.

- All model configuration should be done in a *single* file (i.e., some type of application configuration file). Using this application configuration file it should be possible to reproduce a run of a particular model.

1.1.2. *Scalable...* Aggregate behavior of many (most?) real world social systems fundamentally depends on system size. Therefore in order to accurately model system dynamics we need to be able to run our models as close to observed scale as possible. We need to start thinking about building models “as if” we were developing web applications so that we can leverage recent hardware (i.e., massively multi-core servers, cloud computing architectures that are driving the proliferation of large-scale, distributed web applications in private industry).

1.1.3. *Data-driven...* We want to build models that can be validated against empirical data and therefore we need to design our toolkit to facilitate this. Again we should think of our models “as if” they are web applications. Web applications consume streams of input data and produce streams of output data; some output data streams are processed and then used again as new input data streams. With all of this data flowing around, database integration and data processing are crucial components of web application *design*. We should design our toolkit to integrate cleanly with database architectures, such as [Apache Cassandra](#) and [Neo4j](#), and data processing engines, such as [Apache Spark](#), that are specifically designed to support large-scale web applications.

1.1.4. *Economic...* Finally, we are not seeking to build a general purpose toolkit for agent-based modeling. Rather we want to build a toolkit that is designed to facilitate the construction of agent-based models of economic systems.

An economy is populated with many seemingly disparate types of agents (i.e., consumers, producers, financiers, government, some markets, etc). We need to distill the core essence (in terms of data and behaviors) of these different agents into a multi-layered Application Programming Interface (API) defining a generic *economic agent* that can then be specialized to the various types of economic agents needed for any particular model. Our hope is that by doing this we can reduce the effort needed for the parts of agent-based modeling that consume a great deal of software development time, such as accounting or contract enforcement. We also hope that by identifying the key components we can build standard, highly modular interfaces that make it easy to interchange components of models.

1.1.5. *Reproducible...* Results of many (most?) ABMs are not easily reproducible. More often than not, source code for models used in published articles is not publicly available. Reproducibility is further hindered by the lack of use of “best practices” for software development (in particular unit testing). Documentation is typically lacking.

In order to address these issues and make models built using our toolkit reproducible we should develop the toolkit *from the beginning* as an open-source project and should exemplify “best practices” in software engineering: version control using [Git](#) and [GitHub](#); continuous integration of unit tests using [Travis CI](#) or [Jenkins](#); static analysis of source code using [Codacy](#), [Code Climate](#), etc; documentation should be extensive, written along with the code, and publicly hosted on GitHub.

## 2. HIGH-LEVEL DESIGN

**2.1. Platform architecture.** The high-level design of our platform’s architecture should mimic the layered architecture of a [reactive](#), “Fast Data” web application. Figure 2.1 summarizes the high-level design of our platform architecture.<sup>2</sup>

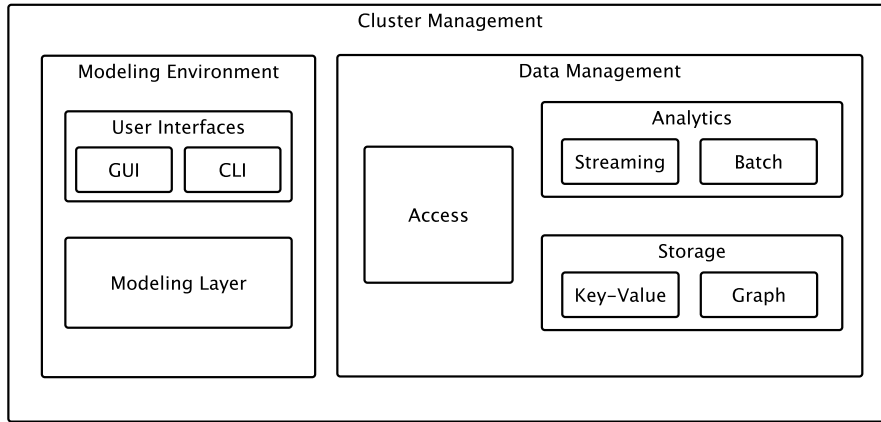


FIGURE 1. High-level architecture design for ScalABM.

**2.1.1. Cluster management.** Our strategy for building and running models at scale is to leverage massively multi-core cloud computing clusters. In order to pursue this scaling strategy our platform will need to have a cluster manager. From Wikipedia: A cluster manager usually is a backend graphical user interface (GUI) or command-line software that runs on one or all cluster nodes (in some cases it runs on a different server or cluster of management servers.) The cluster manager works together with a cluster management agent. These agents run on each node of the cluster to manage and configure services, a set of services, or to manage and configure the complete cluster server itself (see super computing.) In some cases the cluster manager is mostly used to dispatch work for the cluster (or cloud) to perform. In this last case a subset of the cluster manager can be a remote desktop application

<sup>2</sup> See figure 1 from [Wampler \(2015\)](#) for a similar high-level summary of a Reactive, “Fast Data” web application architecture.

that is used not for configuration but just to send work and get back work results from a cluster.

**2.1.2. Modeling environment.** The modeling environment consists of user interfaces and the modeling layer. The “front end” of our modeling environment should consist of two, complementary user interfaces.

- (1) A user-friendly, web-browser based graphical user interface (GUI). The GUI should facilitate interactive exploration of an existing model in near real-time. The GUI should support real-time data streaming, analysis, and visualization.
- (2) An intuitive and consistent command line interface (CLI). In addition to facilitating efficient batch processing of model simulations (i.e., parameter sweeps), the CLI should allow for easy replication of any particular model simulation.

The “back end” of our modeling environment is the modeling layer which consists of the actual source code libraries used to implement our ABMs. The modeling layer should facilitate the construction of new ABMs out of pre-existing, modular components. Model configuration, including specification of all model parameters as well as the “wiring” of model components, should be specified in configuration files that are separate from the actual source code.

**2.1.3. Data Management.** We want to build data-driven ABMs of economic systems. In order for our ABMs to be data-driven, we should design our platform to incorporate data management. There are three components (at least!) to data management: access, analytics, and storage.

- Access: A running ABM will generate a large volume of data some of which will need to be routed for storage, some of which will need to be routed to a data analytics engine, and some of which will need to be routed to log files. Additionally, data will likely flow in the reverse direction. In particular, agents in a running ABM may need to read data from a data store (for example, we might wish to initialize economic agents using real-world data).
- Analytics: A running ABM is a continuous source of data whose volume is not predetermined. Put another way: ABMs generate **reactive data streams**. Our data analytics components should therefore include tooling for processing and analyzing streaming data. In addition to processing and analyzing streaming data, we will also need to perform various “batch” or “mini-batch” computations. Such batch processing jobs would be performed either relatively infrequently on streaming data or upon completion of a model simulation. Our data analytics should include tooling for dealing with batch computations.
- Storage: The way in which we store our data should be heavily influenced by the types of questions we intend to ask of our data.

Broadly speaking, we will need to deal with two types of data: real-world data and model generated data.

There are two kinds of questions that we will want to ask of our model generated data:

- (1) Key-Value databases are designed for storing data in a schema-less way. In a key-value store, each datum consists of an indexed key and a value, hence the name.
- (2) Relationship questions: What does the network structure of interbank lending look like? What does the input-output (IO) production network look like? Graph databases are designed for data whose relations are well represented as a graph and has elements which are interconnected, with an undetermined number of relations between them.

The modeling layer should have read/write access to a scalable data store. Additionally, the data analytics components will need a source of input data. In order to avoid simulations being I/O bound, our data store should have extremely fast write access.

**2.2. Modeling Layer.** Important characteristics of our modeling layer. Scalability. We want to build models with

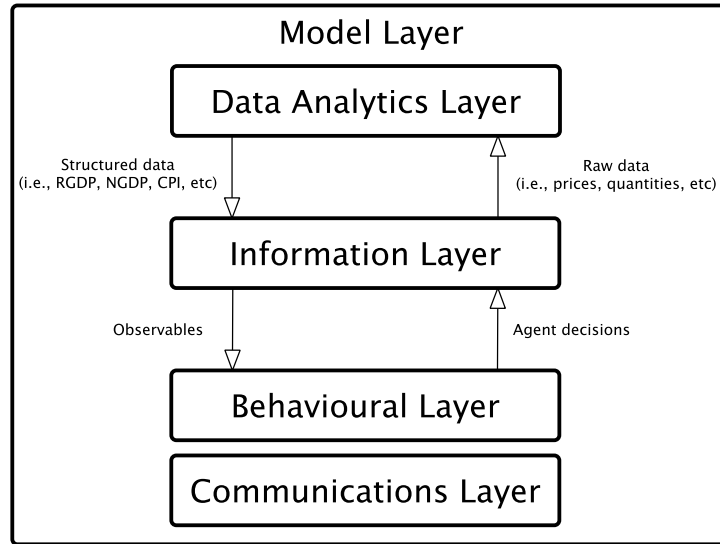


FIGURE 2. High-level design of the modeling layer.

**2.2.1. Communication Layer.** Concurrent communication between real-world economic agents is a fundamental fact of economic life. Inter-agent communication can be either direct (i.e., “peer-to-peer”) or indirect (i.e., via market institutions). In order to model both direct and indirect communication between economic agents, our framework the **Actor model** of concurrency to

model both direct and indirect communication between economic agents as concurrent, asynchronous message passing.

t treats “actors” as the universal primitives of concurrent computation: in response to a message that it receives, an actor can make local decisions, create more actors, send more messages, and determine how to respond to the next message received.

In order to make progress we need to impose some structure on the type of messages passed between agents. This structure will have two forms

Communication protocols: Just as the ACL places structure on the actual messages being passed between agent, the communication protocols impose structure on sequences of messages.

**2.2.2. Behavioral layer.** An economy is populated with many seemingly disparate types of agents (i.e., consumers, producers, financiers, government, etc). Our goal with the behavioral layer is to distill the core essence (in terms of data and behaviors) of these different agents into a multi-layered API defining a generic *economic agent* that can then be specialized to the various types of economic agents needed for any particular model. Note that the API for our economic agent will build upon the communicating agent API.

Several key features distinguish *economic* agents from more generic types of agents. At a minimum these features include: purpose driven (or goal oriented) behaviour, an ability to learn, and the ability to anticipate future events. This suggests that our behavioral layer will need APIs for...

- Goals or objectives (and their associated behavioral rules): Our goals/objectives API needs to be as un-opinionated as possible as prospective users of ScalABM are likely to have strong opinions on how to define appropriate goals and decision rules for their agents. At the same time, we will need to have some type of underlying null model of agent goals/objectives. Utility maximization, Belief-Desire-Intent (BDI), probabilistic discrete choice (AKA, random utility maximization) are possibilities.
- Learning rules: A large number of various learning rules/mechanisms have been proposed in the academic literature. Roughly, learning rules seem to fall into two camps: learning through previous experience (i.e., reinforcement learning) and learning through observation (i.e., belief learning). Useful references for learning rules for ABMs are the two handbook chapters [Brenner \(2006\)](#) and [Duffy \(2006\)](#).
- Expectations formation rules: A large number of various expectation formation rules have been proposed in the literature. Useful references for expectation formation rules are [Hommes \(2006\)](#), [Anufriev and Hommes \(2012\)](#), [Woodford \(2013\)](#), [Assenza et al \(2014\)](#).

2.2.3. *Information Layer.* Our economic agents will need to condition their decisions on an information set. In contrast to (most) DSGE models, information sets in our framework will be highly heterogeneous. The information set for any particular economic agent should consist of three components:

- Private information: information that is idiosyncratic to a particular agent.
- Public information: information that a particular agent shares with one or more agents.
- Global information: information that is shared between all agents.




FIGURE 3. Information set for a particular economic agent.

2.2.4. *Data Analytics Layer.* The data analytics layer encapsulates the functionality for aggregating raw data into additional structured data that agents can use for making decisions. Specifically, the data analytics layer would contain classes and methods for computing standard macro economic indicators (i.e., real and nominal GDP, stock market index, house price index, consumer price index, unemployment rates, etc). Users should be able to extend the data analytics capabilities as needed for a particular model.

### 3. IMPLEMENTATION

#### 3.1. Platform architecture.

3.1.1. *Cluster Management.* We will not need to implement the cluster management layer. Instead we will simply “containerize” an ABM built using our framework and send it off to a third-party cloud computing service provider such as [AWS](#), [Heroku](#), [Mesosphere](#), etc. We can then choose the scale at which we wish to run the model. Once the model is running on the cluster we can configure the scale(s) at which we wish to simulate the model and the cloud computing service provider will handle the rest.

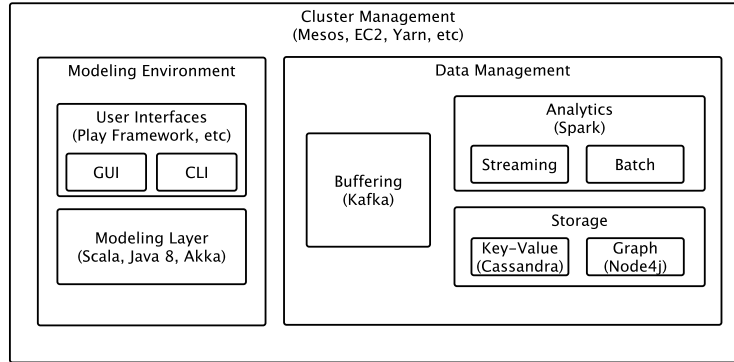


FIGURE 4. Macro-grain architecture for ScalABM.

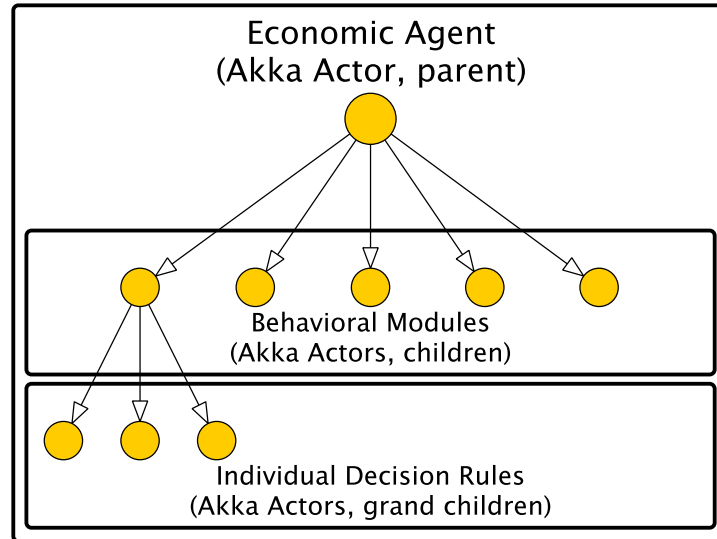


FIGURE 5. ???

### 3.2. Modeling Layer.

**3.2.1. Communication Layer.** In order to impose structure on inter-agent communications we have developed an API for a scalable agent communication language for economic agents tentatively called ScalACL. ScalACL leverages the **Actor model** of concurrency as implemented in the **Akka** library to model both direct and indirect communication between economic agents using concurrent, asynchronous message passing.<sup>3</sup> The ScalACL API specifies...

<sup>3</sup> Akka is a toolkit and runtime for building highly concurrent, distributed, and resilient message-driven applications on the JVM.



- a set of abstract message types that impose structure on the messages passed between a group of economic agents,
- a set of abstract protocols that impose structure on conversations (i.e., sequences of messages) between a group of economic agents,
- a behavioral trait that allows an agent to communicate using the language.

Each abstract message type can be thought of as defining an “envelope” containing the actual content of the message that is to be exchanged between a group of economic agents. Defining envelopes containing messages is useful because it allows agents to react based on the type message received. Each abstract protocol defines a particular subset of message types that can be sent by an agent in response to a particular type of message that it has received.

Our agent communication API is influenced by, but not slave to, the [Foundation for Intelligent Physical Agents \(FIPA\)](#) compliant [Agent Communication Language \(ACL\)](#).